



Multilayer perceptron

Backpropagation and Neural Processing




Outline



- **Multi-layer perceptron**
 - non-linearity
 - gradient descent
- **Propagation**
- **Backpropagation**
 - network updates
 - tensor operations
 - batch *versus* stochastic updates
 - cross-entropy
- **Learning convergence**
 - optimality
 - early stopping



Outline



- **Multi-layer perceptron**
 - **non-linearity**
 - gradient descent
- Propagation
- Backpropagation
 - network updates
 - tensor operations
 - batch *versus* stochastic updates
 - cross-entropy
- Learning convergence
 - optimality
 - early stopping


Perceptron limitations

- Designed for binary classification $|Z| = 2$
 - how can we extend for other predictive tasks?
- Linear separation only (the XOR problem)




Non-linear problems

- In mathematics, we can compose **linear** functions, $f(x)$ and $g(x)$, to form a **non-linear** function, $h(f(x), g(x))$, to model more complex behavior
- What if we combine multiple perceptrons? Are we able to solve the XOR problem?
 - YES!
- We can create an intermediate or hidden layer of perceptrons
 - a network with a single hidden unit can already represent any Boolean function!



Multi-layer networks


- The composition of perceptrons can be **organized in layers**
 - the outputs of perceptrons in one layer feed the input of perceptrons in the next layer
 - hence **multi-layer perceptrons** (MLPs) are also termed **feed forward networks**
 - a simplified architecture to facilitate the learning
 - the presence of cycles (e.g. two perceptrons feeding each other) brings complexities
 - the first and last layers are referred as input and output layers



Non-linear models

- Yet... multiple layers of cascade linear units still produce only linear functions

$$z_{out} = \sum w \left(\sum w x_i \right)$$




– solution?


$$z_{out} = f \left(\sum w f \left(\sum w x_i \right) \right)$$

Activation functions





- networks able of representing nonlinear functions:
 - use nonlinear activation functions
 - activations should be continuous and differentiable
 - problems of sign and step?



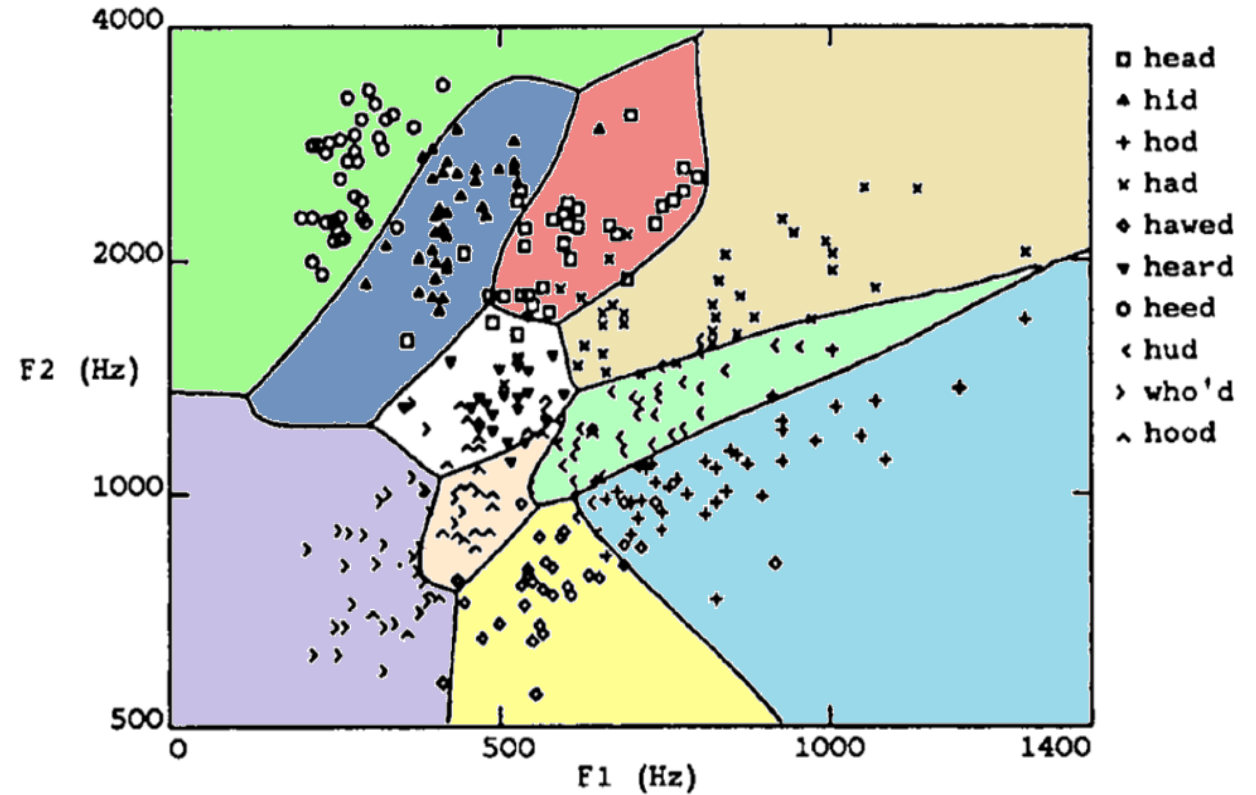

Sign Function




Step Function



Non-linear boundaries




Solving the XOR problem



- *Exercise:* place the necessary activation functions on the perceptrons to yield the desired result

MLPs for regression and classification


- Until now... a single perceptron can only answer a binary classification task
- Neural networks can be extended to handle more general predictive tasks
 - *multiclass classification*
 - **regression**
 - replace the last sign/sigmoid/tanh activation by other function – e.g. [rectified] linear unit
 - **multiple-output prediction**
 - multiple targets
(e.g., autonomous driving – speed and direction)




- What about the learning? The same principles apply as we will see later

Multi-class classification

- Many classification tasks have high cardinality
 - e.g. document categorization, product recommendation, character recognition
- Given a set L of class labels
 - $|L|$ output nodes, one per class
 - predicted class is the output neuron with the higher value




Outline



- **Multi-layer perceptron**
 - non-linearity
 - **gradient descent**
- Propagation
- Backpropagation
 - network updates
 - tensor operations
 - batch *versus* stochastic updates
 - cross-entropy
- Learning convergence
 - optimality
 - early stopping

Learning multi-layer networks

- The great power of multi-layer networks was realized a long ago
 - yet only in the 80s it was shown how to make them learn!
- **Backpropagation** is the most used learning algorithm for multi-layer neural networks
 - invented independently several times
 - Bryson and Ho [1969]
 - Werbos [1974]
 - Parker [1985]
 - Rumelhart et al. [1986]
 - See Parallel Distributed Processing - Vol. 1, Foundations



Backpropagation

- **Goal**

- given a training set of input-output pairs $\{\mathbf{x}_i, t_i\}$
- learn the parameters of the network, w_{ij}

- **How?**

- finding the weights that minimize the loss
 - back to the usual error function to assess our network

$$E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n (t^{(i)} - o^{(i)})^2$$

- for L outputs and n input-output pairs $\{\mathbf{x}_i, \mathbf{t}_i\}$


$$E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n \sum_{l=1}^L (t_l^{(i)} - o_l^{(i)})^2$$

Backpropagation


- How do we find the weights that minimize the given loss?
 - gradient descent! Adjust the weights in the direction where the error decreases

$$\Delta w_{ij} = -\eta \frac{\partial E(\mathbf{w})}{\partial w_{ij}}$$

- two major steps:
 - *forward propagation* of inputs until the end of the network
 - compute the observed errors δ and propagate them *backwards* (hence **backpropagation**)



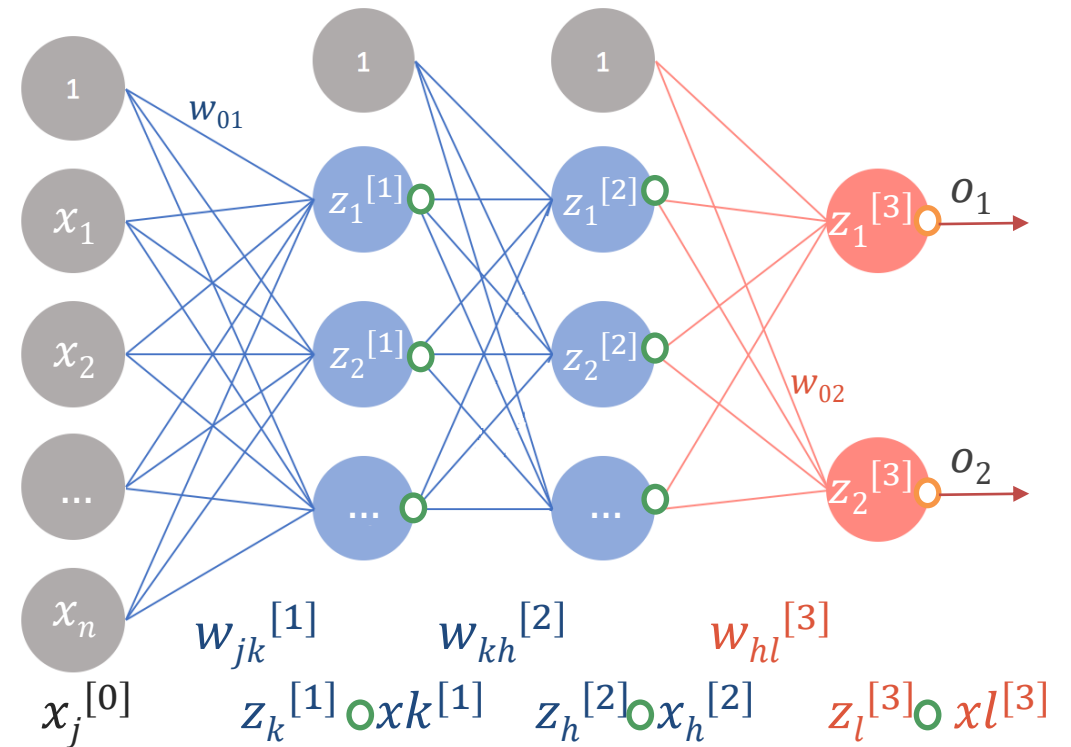
Outline



- Multi-layer perceptron
 - non-linearity
 - task formulations
- **Propagation**
- Backpropagation
 - gradient descent updates
 - tensor operations
 - batch *versus* stochastic updates
 - cross-entropy
- Learning convergence
 - optimality
 - early stopping

Notation

- network with P **layers**, $p = 1, 2, \dots, P$
- $z_k^{[p]}$ is the **net value** of the k^{th} node of the p^{th} layer
- $\phi^{[p]}$ is the activation function of the p^{th} layer
- $x_k^{[p]}$ is the **signal** of the k^{th} node of the p^{th} layer
 - $x_k^{[0]}$ is a synonym for x_k , i.e. the k^{th} **input feature**
 - $x_k^{[P]}$ is a synonym for o_k , i.e. the k^{th} **network output**
- $w_{ij}^{[p]}$ is the **weight** of the connection from i^{th} node of the $p-1^{th}$ layer to the j^{th} node of the p^{th} layer
- **bias** is the product of $x_0^{[p]} = 1$ and $w_{0k}^{[p]}$
 - similarly to the perceptron and linear regression, we consider a bias term to aid the learning
 - the bias is present on *every* layer




Propagation

- Let us consider a simple three-layered network
- Given the observation \mathbf{x} the *net* values of the first layer

$$net_k^{[1]} = z_k^{[1]} = \sum_{j=1}^m w_{jk}^{[1]} x_j^{[0]} = \sum_{j=1}^m w_{jk}^{[1]} x_j$$

– ... and produces the **output**

$$x_k^{[1]} = f(z_k^{[1]}) = \phi^{[1]} \left(\sum_{j=1}^m w_{jk}^{[1]} x_j \right)$$



Propagation


– similarly considering the second layer...

$$x_h^{[2]} = \phi^{[2]}(z_h^{[2]}) = \phi^{[2]} \left(\sum_{k=1}^{m^{[1]}} w_{kh}^{[2]} x_k^{[1]} \right)$$

$$= \phi^{[2]} \left(\sum_{k=1}^{m^{[1]}} w_{kh}^{[2]} \phi^{[1]} \left(\sum_{j=1}^m w_{jk}^{[1]} x_j \right) \right)$$

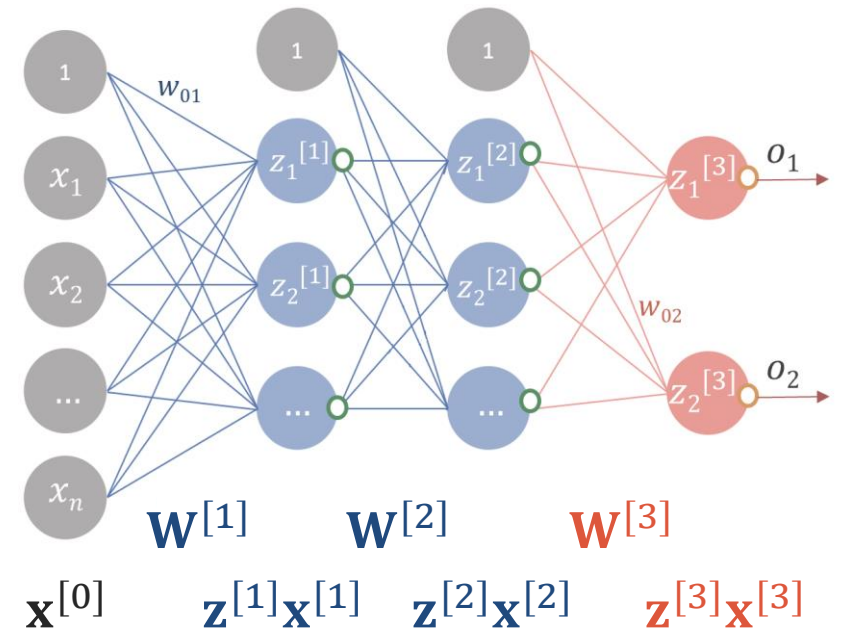
– and the **output layer**...

$$o_l = x_l^{[3]} = \phi^{[3]}(z_l^{[3]}) = \phi^{[3]} \left(\sum_{h=1}^{m^{[2]}} w_{hl}^{[3]} x_h^{[2]} \right) = \phi^{[3]} \left(\sum_{h=1}^{m^{[2]}} w_{hl}^{[3]} (\dots) \right)$$



Vector notation

- network with P layers, $p = 1, 2, \dots, P$
- **parameters:**
 - $\mathbf{W}^{[p]}$ contains the **weights** connecting $p-1^{th}$ layer and p^{th} layer
 - either $\mathbf{W}^{[p]}$ contains all weights or **biases** are isolated in $\mathbf{b}^{[p]}$
- **variables:**
 - $\mathbf{z}^{[p]}$ are the **net values** of the p^{th} layer
 - $\mathbf{x}^{[p]}$ is the activated/output values of the p^{th} layer, $\mathbf{x}^{[p]} = \phi^{[1]}(\mathbf{z}^{[p]})$
 - $\mathbf{x}^{[0]}$ is a synonym for \mathbf{x} (input **observation**)
 - $\mathbf{x}^{[P]}$ is a synonym for \mathbf{o} , i.e. the network **output**



Propagation

- Consider the given network, assuming $m = 4$ and hidden layers with 3 nodes each
- The **parameters** of the given neural network:

$$\begin{aligned} - \mathbf{W}^{[1]} &= \begin{pmatrix} w_{11}^{[1]} & \dots & w_{41}^{[1]} \\ \dots & & \dots \\ w_{13}^{[1]} & \dots & w_{43}^{[1]} \end{pmatrix}, \mathbf{W}^{[2]} = \begin{pmatrix} w_{11}^{[2]} & \dots & w_{31}^{[2]} \\ \dots & & \dots \\ w_{13}^{[2]} & \dots & w_{33}^{[2]} \end{pmatrix}, \mathbf{W}^{[3]} = \begin{pmatrix} w_{11}^{[3]} & w_{21}^{[3]} & w_{31}^{[3]} \\ w_{12}^{[3]} & w_{22}^{[3]} & w_{32}^{[3]} \end{pmatrix} \\ - \mathbf{b}^{[1]} &= \begin{pmatrix} b_1^{[1]} \\ \dots \\ b_3^{[1]} \end{pmatrix}, \mathbf{b}^{[2]} = \begin{pmatrix} b_1^{[2]} \\ \dots \\ b_3^{[2]} \end{pmatrix}, \mathbf{b}^{[3]} = \begin{pmatrix} b_1^{[3]} \\ b_2^{[3]} \end{pmatrix} \end{aligned}$$

- The **nodes** of the neural network

$$\begin{aligned} - \mathbf{x} = \mathbf{x}^{[0]} &= \begin{pmatrix} x_1 \\ \dots \\ x_4 \end{pmatrix}, \mathbf{x}^{[1]} = \phi^{[1]}(\mathbf{W}^{[1]}\mathbf{x} + \mathbf{b}^{[1]}) = \begin{pmatrix} x_1^{[1]} \\ \dots \\ x_3^{[1]} \end{pmatrix}, \mathbf{x}^{[2]} = \phi^{[2]}(\mathbf{W}^{[2]}\mathbf{x}^{[1]} + \mathbf{b}^{[2]}) = \begin{pmatrix} x_1^{[2]} \\ \dots \\ x_3^{[2]} \end{pmatrix} \\ - \mathbf{o} = \mathbf{x}^{[3]} &= \phi^{[3]}(\mathbf{W}^{[3]}\mathbf{x}^{[2]} + \mathbf{b}^{[3]}) = \begin{pmatrix} o_1 \\ o_2 \end{pmatrix} \end{aligned}$$

Propagation: example

- Considering all weights initialized at 0.1, no biases, and sigmoid σ activation function

What is the output for observation $\mathbf{x}^T = (1 \ 1 \ 0 \ 0)$?


$$- \mathbf{W}^{[1]} = \begin{pmatrix} 0.1 & 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 & 0.1 \end{pmatrix}, \mathbf{W}^{[2]} = \begin{pmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{pmatrix}, \mathbf{W}^{[3]} = \begin{pmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{pmatrix}, \mathbf{b}^{[1]} = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}, \mathbf{b}^{[2]} = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}, \mathbf{b}^{[3]} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

$$- \mathbf{x}^{[1]} = \phi^{[1]}(\mathbf{W}^{[1]}\mathbf{x} + \mathbf{b}^{[1]}) = \sigma \left(\begin{pmatrix} 0.1 & 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 & 0.1 \end{pmatrix} \begin{pmatrix} 1 \\ 1 \\ 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} \right) = \sigma \begin{pmatrix} 0.2 \\ 0.2 \\ 0.2 \end{pmatrix},$$

$$- \mathbf{x}^{[2]} = \phi^{[2]}(\mathbf{W}^{[2]}\mathbf{x}^{[1]} + \mathbf{b}^{[2]}) = \sigma \left(\begin{pmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{pmatrix} \begin{pmatrix} \sigma(0.2) \\ \sigma(0.2) \\ \sigma(0.2) \end{pmatrix} + \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} \right) = \sigma \begin{pmatrix} 0.165 \\ 0.165 \\ 0.165 \end{pmatrix}$$

$$- \mathbf{o} = \mathbf{x}^{[3]} = \phi^{[3]}(\mathbf{W}^{[3]}\mathbf{x}^{[2]} + \mathbf{b}^{[3]}) = \sigma \left(\begin{pmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{pmatrix} \begin{pmatrix} \sigma(0.165) \\ \sigma(0.165) \\ \sigma(0.165) \end{pmatrix} + \begin{pmatrix} 0 \\ 0 \end{pmatrix} \right) = \sigma \begin{pmatrix} 0.162 \\ 0.162 \end{pmatrix} = \begin{pmatrix} 0.54 \\ 0.54 \end{pmatrix}$$

Outline



- Multi-layer perceptron
 - non-linearity
 - gradient descent
- Propagation
- **Backpropagation**
 - **network updates**
 - tensor operations
 - batch *versus* stochastic updates
 - cross-entropy
- Learning convergence
 - optimality
 - early stopping

Loss: sum of squared errors

- Recall our goal: learn the parameters of the network
 - optimize weights \mathbf{w} – matrices $\mathbf{W}^{[p]}$ and biases $\mathbf{b}^{[p]}$ for $p \in \{1, 2, \dots, P\}$
- **How?**
 - minimize error $E(\mathbf{w})$ to estimate \mathbf{w} ! Back to our sum of squared errors (SSE)

$$E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n (t^{(i)} - o^{(i)})^2$$

- for L outputs and n input-output pairs $\{\mathbf{x}_i, \mathbf{t}_i\}$

$$E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n \sum_{l=1}^L (t_l^{(i)} - o_l^{(i)})^2$$

- for our previous example:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n \sum_{l=1}^2 \left(t_l^{(i)} - \phi^{[3]} \left(\sum_{h=1}^3 w_{hl}^{[3]} \cdot \phi^{[2]} \left(\sum_{k=1}^3 w_{kh}^{[2]} \phi^{[1]} \left(\sum_{j=1}^m w_{jk}^{[1]} x_j^{(i)} \right) \right) \right) \right)^2$$

Backpropagation

- $E(\mathbf{w})$ is differentiable if $\phi^{[p]}$ are differentiable
 - **Gradient Descent** can be applied!
 - similarly to what we did for the perceptron, we can infer the update rules for a MLP
- Yet before advancing a central tool – **chain rule** – to differentiate composite functions:

$$\frac{\partial f(g(x))}{\partial x} = f'(g(x)) \frac{\partial g(x)}{\partial x} = f'(g(x))g'(x)$$


$$\nabla_{\mathbf{x}} f(g(\mathbf{x})) = \frac{\partial f(g(\mathbf{x}))}{\partial \mathbf{x}} = \frac{\partial f(g(\mathbf{x}))}{\partial g} \cdot \frac{\partial g(\mathbf{x})}{\partial \mathbf{x}}$$

– example:

$$\frac{\partial \exp(3x^2)}{\partial x} = \frac{\exp(f(x))}{\partial f} \frac{\partial (3x^2)}{\partial x} = \exp(3x^2) \times 6x$$

Backpropagation

- To learn a MLP we iteratively select observations and:
 - (forward) propagation of observation values
 - $\mathbf{x}^{[p]} = \phi^{[p]}(\mathbf{W}^{[p]}\mathbf{x}^{[p-1]} + \mathbf{b}^{[p]})$
 - backpropagation the errors
 - compute $\delta^{[p]}$ differences to expectations from last to first layer
 - update the weights
 - $\mathbf{W}^{[p]} = \mathbf{W}^{[p]} - \eta \frac{\partial E}{\partial \mathbf{W}^{[p]}}$ where $\frac{\partial E}{\partial \mathbf{W}^{[p]}} = \delta^{[p]}(\mathbf{x}^{[p-1]})^T$
 - $\mathbf{b}^{[p]} = \mathbf{b}^{[p]} - \eta \frac{\partial E}{\partial \mathbf{b}^{[p]}}$ where $\frac{\partial E}{\partial \mathbf{b}^{[p]}} = \delta^{[p]}$




Backpropagation

- To compute $\delta^{[p]}$, we need first to infer the update rule...
- Let us infer the update rules on the right
 - let us start with the **last layer**...

$$\begin{aligned}\Delta w_{kl}^{[2]} &= -\eta \frac{\partial E}{\partial w_{kl}^{[2]}} = -\eta \frac{\partial}{\partial w_{kl}^{[2]}} \sum_{i=1}^n \left(t_l^{(i)} - o_l^{(i)}\right)^2 \\ &= -\eta \sum_{i=1}^n \left(t_l^{(i)} - o_l^{(i)}\right) \cdot \left(-\phi'^{[2]}(z_l^{(i)[2]})\right) \cdot x_k^{(i)[2]} \\ &= \eta \sum_{i=1}^n \delta_l^{(i)[2]} \cdot x_k^{(i)[2]}\end{aligned}$$

where $\delta_l^{[2]} = (t_l - o_l)\phi'^{[2]}(z_l^{[2]})$



Backpropagation

- ... let us continue with the **hidden layer**. Using the chain rule:

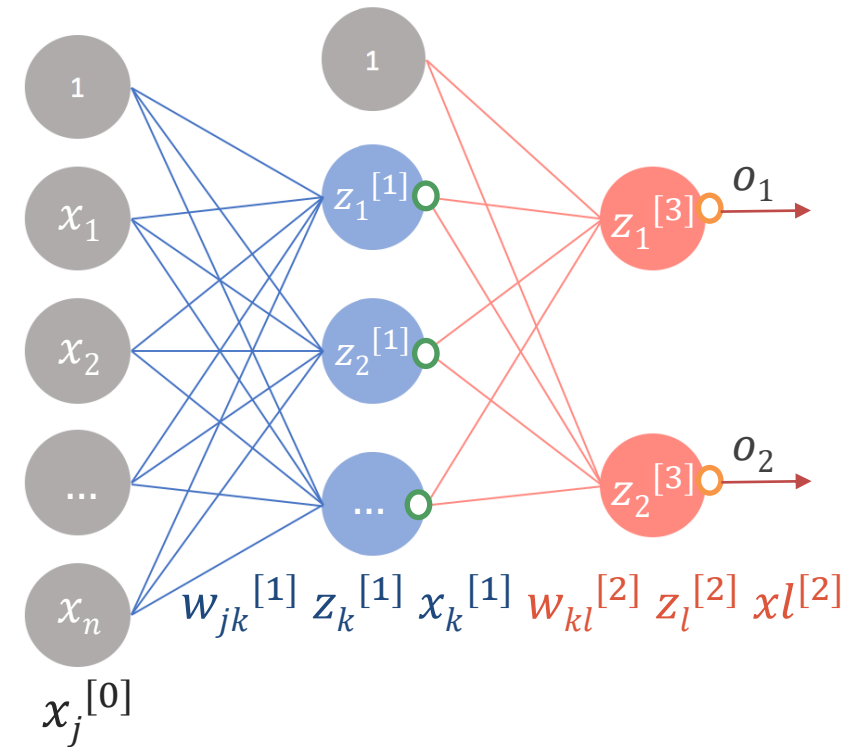
$$\Delta w_{jk}^{[1]} = -\eta \frac{\partial E}{\partial w_{jk}^{[1]}} = -\eta \sum_{i=1}^n \frac{\partial E}{\partial x_k^{(i)[1]}} \cdot \frac{\partial x_k^{(i)[1]}}{\partial w_{jk}^{[1]}}$$

$$\frac{\partial x_k^{(i)[1]}}{\partial w_{jk}^{[1]}} = \frac{\partial}{\partial w_{jk}^{[1]}} \left(\phi^{[1]} \left(\sum_{j=1}^4 w_{jk}^{[1]} x_j^{(i)} \right) \right) = \phi'^{[1]}(z_k^{(i)[1]}) \cdot x_j^{(i)}$$

$$\frac{\partial E}{\partial x_k^{(i)[1]}} = \sum_{l=1}^2 \delta_l^{(i)[2]} w_{kl}^{(i)[2]}$$

$$\Delta w_{jk}^{[1]} = \eta \sum_{i=1}^n \delta_k^{(i)[1]} x_j^{(i)}$$

where $\delta_k^{[1]} = \phi'^{[1]}(z_k^{[1]}) \sum_{l=1}^2 \delta_l^{[2]} w_{kl}^{[2]}$



Backpropagation

- In general, with an **arbitrary number of layers**, the back-propagation **update rule** has the form

$$\Delta w_{jk}^{[p]} = \eta \sum_{i=1}^n \delta_k^{(i)[p]} x_j^{(i)[p-1]}$$

- where the δ of the last layer is given by

$$\delta_k^{[P]} = (t_k - o_k) \phi'^{[P]}(z_k^{[P]})$$

- and the remaining δ given by

$$\delta_k^{[p]} = \phi'^{[p]}(z_k^{[p]}) \sum_{h=1}^{m^{[p+1]}} \delta_h^{[p+1]} w_{kh}^{[p+1]}$$


- The deltas essentially propagate the difference between expected and observed outputs with correction factors – defining the strength to change weights

Activation functions

- we have to use a *differentiable* activation functions $\phi^{[p]}$
 - sigmoid, hyperbolic tangent and rectified linear unit (ReLU) are common options
 - we can use different activations for different layers (e.g. $\phi^{[p]} \neq \phi^{[p+1]}$)
 - yet the nodes within the same layer generally have the same activation function


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Outline



- Multi-layer perceptron
 - non-linearity
 - gradient descent
- Propagation
- **Backpropagation**
 - network updates
 - **tensor operations**
 - batch *versus* stochastic updates
 - cross-entropy
- Learning convergence
 - optimality
 - early stopping

Vector notation

- Let us recover important matricial operations
 - **Hadamard product**: element-wise product of vectors or matrices

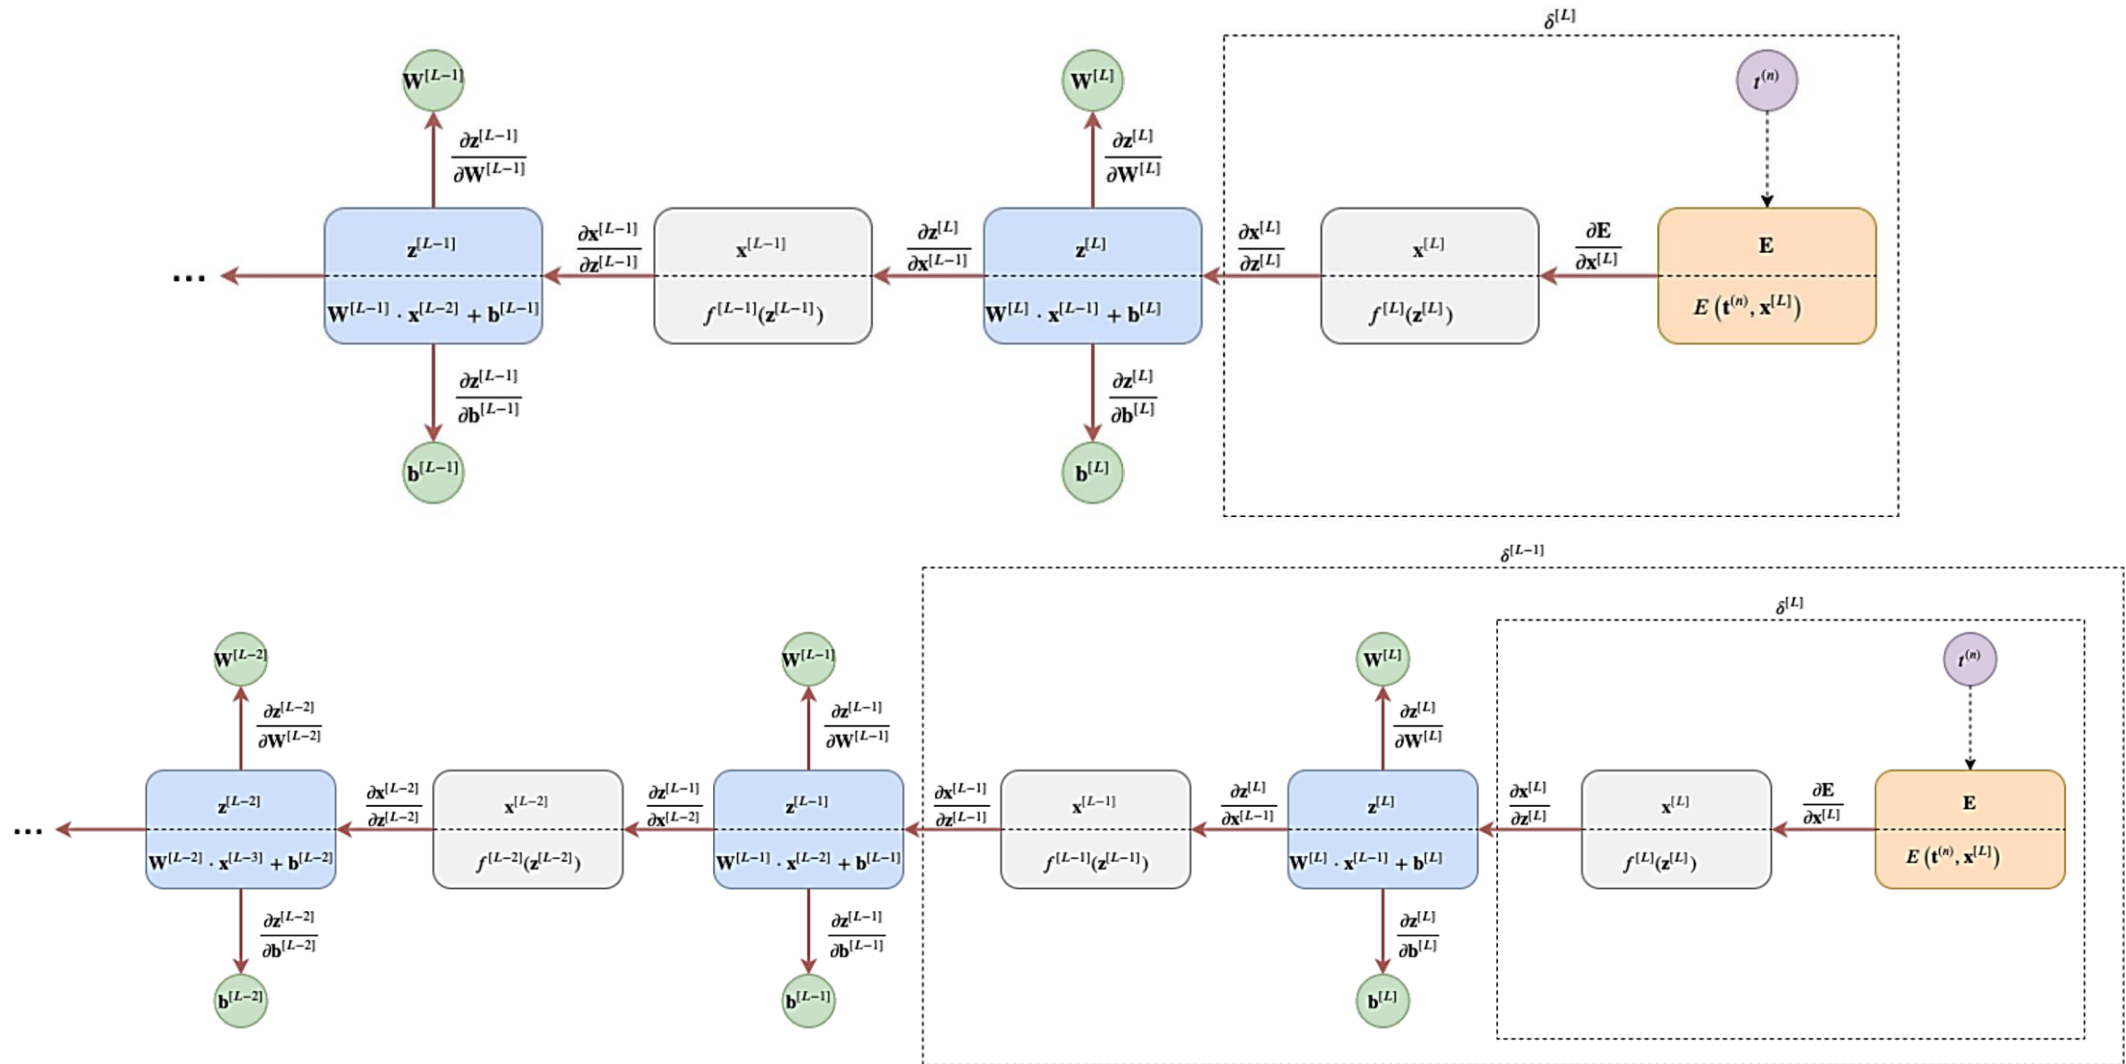
$$\begin{pmatrix} a_1 & a_2 \\ a_3 & a_4 \end{pmatrix} \circ \begin{pmatrix} b_1 & b_2 \\ b_3 & b_4 \end{pmatrix} = \begin{pmatrix} a_1 b_1 & a_2 b_2 \\ a_3 b_3 & a_4 b_4 \end{pmatrix}$$

- **tensor product**: outer product

$$\begin{pmatrix} \omega_0 \\ \omega_1 \end{pmatrix} \otimes \begin{pmatrix} \omega_0 \\ \omega_1 \end{pmatrix} = \begin{pmatrix} \omega_0 \cdot \omega_0 \\ \omega_0 \cdot \omega_1 \\ \omega_1 \cdot \omega_0 \\ \omega_1 \cdot \omega_1 \end{pmatrix}$$

$$A \otimes B = \begin{pmatrix} a_{11} \cdot B & a_{12} \cdot B \\ a_{21} \cdot B & a_{22} \cdot B \end{pmatrix} = \begin{pmatrix} a_{11} \cdot b_{11} & a_{11} \cdot b_{12} & a_{12} \cdot b_{11} & a_{12} \cdot b_{12} \\ a_{11} \cdot b_{21} & a_{11} \cdot b_{22} & a_{12} \cdot b_{21} & a_{12} \cdot b_{22} \\ a_{21} \cdot b_{11} & a_{21} \cdot b_{12} & a_{22} \cdot b_{11} & a_{22} \cdot b_{12} \\ a_{21} \cdot b_{21} & a_{21} \cdot b_{22} & a_{22} \cdot b_{21} & a_{22} \cdot b_{22} \end{pmatrix}$$

Backpropagation in vector notation



Backpropagation

- Let us now rewrite the updating rules with gradient descent using vector notation
 - let us minimize the error: $E(\mathbf{w}) = \frac{1}{2}(\mathbf{t} - \mathbf{o})^2$
 - applying the chain rule taking into attention matricial operators (see previous slide)...

$$\frac{\partial E}{\partial \mathbf{W}^{[p]}} = \frac{\partial E}{\partial \mathbf{x}^{[p]}} \circ \frac{\partial \mathbf{x}^{[p]}}{\partial \mathbf{z}^{[p]}} \cdot \frac{\partial \mathbf{z}^{[p]}}{\partial \mathbf{W}^{[p]}}$$

$$\frac{\partial \mathbf{x}^{[p]}}{\partial \mathbf{z}^{[p]}} = \phi'^{[p]}(\mathbf{z}^{[p]})$$

$$\frac{\partial \mathbf{z}^{[p]}}{\partial \mathbf{W}^{[p]}} = \mathbf{x}^{[p-1]}$$

$$(\text{for last layer } p = P) \quad \frac{\partial E}{\partial \mathbf{x}^{[P]}} = \frac{1}{2} \frac{\partial}{\partial \mathbf{x}^{[P]}} (\mathbf{t} - \mathbf{x}^{[P]})^2 = \mathbf{x}^{[P]} - \mathbf{t}$$

Backpropagation

- Finally, the update of **weights**:

$$\mathbf{W}^{[p]} = \mathbf{W}^{[p]} - \eta \frac{\partial E}{\partial \mathbf{W}^{[p]}}, \text{ where } \frac{\partial E}{\partial \mathbf{W}^{[p]}} = \boldsymbol{\delta}^{[p]} \frac{\partial \mathbf{z}^{[p]}}{\partial \mathbf{W}^{[p]}} = \boldsymbol{\delta}^{[p]} (\mathbf{x}^{[p-1]})^T$$

- where the delta from the **last layer** is

$$\boldsymbol{\delta}^{[P]} = \frac{\partial E}{\partial \mathbf{x}^{[P]}} \circ \frac{\partial \mathbf{x}^{[P]}}{\partial \mathbf{z}^{[P]}} = (\mathbf{x}^{[P]} - \mathbf{t}) \circ \phi'^{[P]}(\mathbf{z}^{[P]})$$

- and the deltas from the **remaining layers** (using recursion):

$$\boldsymbol{\delta}^{[p]} = \left(\frac{\partial \mathbf{z}^{[p+1]}}{\partial \mathbf{x}^{[p]}} \right)^T \cdot \boldsymbol{\delta}^{[p+1]} \circ \frac{\partial \mathbf{x}^{[p]}}{\partial \mathbf{z}^{[p]}} = \mathbf{W}^{[p+1]T} \cdot \boldsymbol{\delta}^{[p+1]} \circ \phi'^{[p]}(\mathbf{z}^{[p]})$$

Backpropagation

- The update of **biases** is analogous:

$$\mathbf{b}^{[p]} = \mathbf{b}^{[p]} - \eta \frac{\partial E}{\partial \mathbf{b}^{[p]}}, \text{ where } \frac{\partial E}{\partial \mathbf{b}^{[p]}} = \delta^{[p]} \frac{\partial \mathbf{z}^{[p]}}{\partial \mathbf{b}^{[p]}} = \delta^{[p]}$$

– since $\frac{\partial E}{\partial \mathbf{b}^{[p]}} = \frac{\partial E}{\partial \mathbf{x}^{[p]}} \circ \frac{\partial \mathbf{x}^{[p]}}{\partial \mathbf{z}^{[p]}} \cdot \frac{\partial \mathbf{z}^{[p]}}{\partial \mathbf{b}^{[p]}}$ and $\frac{\partial \mathbf{z}^{[p]}}{\partial \mathbf{b}^{[p]}} = \mathbf{1}$

- Recall the full learning process

do following steps **until** convergence (e.g. loss not improving on the last d steps):

- select a set of observations $\{\mathbf{x}_{i_1}, \dots, \mathbf{x}_{i_q}\}$ from X
- updated the parameters $\mathbf{W}^{[p]}, \mathbf{b}^{[p]}$ of the network using the backpropagation rules

Backpropagation: example

- Let us recover our example (network on the right)
 - hidden layers with 3 nodes each, all weights at 0.1, no biases
 - sigmoid activation function $\phi^{[p]}(x) = \sigma(x)$
 - learning rate $\eta = 1$ and squared error loss
 - quest:** update parameters using $\mathbf{x}^T = (1 \ 1 \ 0 \ 0)$ with $\mathbf{t} = (0 \ 1)$

- Solution notes**

- for convenience, let us recall the sigmoid derivative

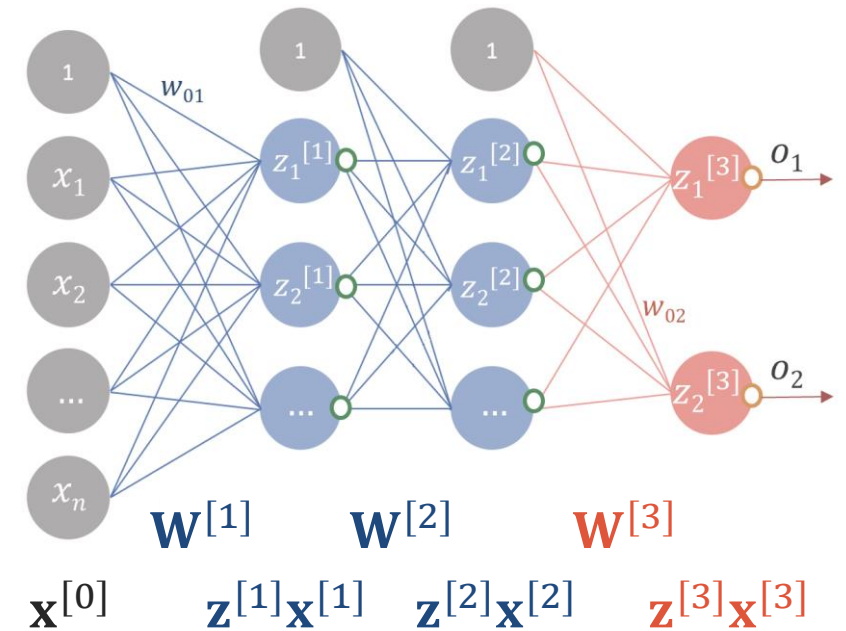
$$\frac{\partial \sigma(x)}{\partial x} = \frac{\partial}{\partial x} \frac{1}{1 + \exp(-x)} = \sigma(x)(1 - \sigma(x))$$

Note: $\sigma(x) = f(g(h(q(x))))$

where $q(x) = -x$, $h(x) = \exp(x)$, $g(x) = 1 + \exp(x)$, $f(x) = \frac{1}{x}$

- first, let us recover the propagated values

$$\mathbf{x}^{[1]} = \sigma \begin{pmatrix} 0.2 \\ 0.2 \\ 0.2 \end{pmatrix}, \mathbf{x}^{[2]} = \sigma \begin{pmatrix} 0.165 \\ 0.165 \\ 0.165 \end{pmatrix}, \mathbf{o} = \mathbf{x}^{[3]} = \sigma \begin{pmatrix} 0.162 \\ 0.162 \end{pmatrix}$$



Backpropagation: example

- Third, let us compute the deltas:

$$- \boldsymbol{\delta}^{[3]} = (\mathbf{x}^{[3]} - \mathbf{t}) \circ \sigma(\mathbf{z}^{[3]}) \circ (1 - \sigma(\mathbf{z}^{[3]})) = \left(\begin{pmatrix} 0.54 \\ 0.54 \end{pmatrix} - \begin{pmatrix} 1 \\ 0 \end{pmatrix} \right) \circ \begin{pmatrix} 0.54 \\ 0.54 \end{pmatrix} \circ \left(1 - \begin{pmatrix} 0.54 \\ 0.54 \end{pmatrix} \right) = \begin{pmatrix} -0.114 \\ 0.134 \end{pmatrix}$$

$$- \boldsymbol{\delta}^{[2]} = \mathbf{W}^{[3]T} \cdot \boldsymbol{\delta}^{[3]} \circ \sigma(\mathbf{z}^{[2]}) \circ (1 - \sigma(\mathbf{z}^{[2]})) = \begin{pmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{pmatrix} \cdot \begin{pmatrix} -0.114 \\ 0.134 \end{pmatrix} \circ \sigma \begin{pmatrix} 0.165 \\ 0.165 \\ 0.165 \end{pmatrix} \left(1 - \sigma \begin{pmatrix} 0.165 \\ 0.165 \\ 0.165 \end{pmatrix} \right) = \begin{pmatrix} 0.0005 \\ 0.0005 \\ 0.0005 \end{pmatrix}$$

$$- \boldsymbol{\delta}^{[1]} = \mathbf{W}^{[2]T} \cdot \boldsymbol{\delta}^{[2]} \circ \sigma(\mathbf{z}^{[1]}) \circ (1 - \sigma(\mathbf{z}^{[1]})) = \begin{pmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{pmatrix} \cdot \begin{pmatrix} 0.0005 \\ 0.0005 \\ 0.0005 \end{pmatrix} \circ \sigma \begin{pmatrix} 0.2 \\ 0.2 \\ 0.2 \end{pmatrix} \left(1 - \sigma \begin{pmatrix} 0.2 \\ 0.2 \\ 0.2 \end{pmatrix} \right) = \begin{pmatrix} 3.7E-5 \\ 3.7E-5 \\ 3.7E-5 \end{pmatrix}$$

- Fourth, and finally, let us perform the updates of the weights...

$$\bullet \mathbf{W}^{[1]} = \mathbf{W}^{[1]} - \eta \frac{\partial E}{\partial \mathbf{W}^{[1]}} = \mathbf{W}^{[1]} - \eta \boldsymbol{\delta}^{[1]} \mathbf{x}^T = \begin{pmatrix} 0.1 & 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 & 0.1 \end{pmatrix} - 1 \begin{pmatrix} 3.7E-5 \\ 3.7E-5 \\ 3.7E-5 \end{pmatrix} \begin{pmatrix} 1 & 1 & 0 & 0 \end{pmatrix} = \begin{pmatrix} .09996 & .09996 & 0.1 & 0.1 \\ .09996 & .09996 & 0.1 & 0.1 \\ .09996 & .09996 & 0.1 & 0.1 \end{pmatrix}$$

$$\bullet \mathbf{b}^{[1]} = \mathbf{b}^{[1]} - \eta \frac{\partial E}{\partial \mathbf{b}^{[1]}} = \mathbf{b}^{[1]} - \eta \boldsymbol{\delta}^{[1]} = \begin{pmatrix} -3.7E-5 \\ -3.7E-5 \\ -3.7E-5 \end{pmatrix}$$

Backpropagation: example

$$\bullet \mathbf{W}^{[2]} = \mathbf{W}^{[2]} - \eta \frac{\partial E}{\partial \mathbf{W}^{[2]}} = \mathbf{W}^{[2]} - \eta \boldsymbol{\delta}^{[2]} (\mathbf{x}^{[1]})^T = \begin{pmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{pmatrix} - 1 \begin{pmatrix} 0.0005 \\ 0.0005 \\ 0.0005 \end{pmatrix} (0.55 \quad 0.55 \quad 0.55) = \begin{pmatrix} 0.0997 & 0.0997 & 0.0997 \\ 0.0997 & 0.0997 & 0.0997 \\ 0.0997 & 0.0997 & 0.0997 \end{pmatrix}$$

$$\bullet \mathbf{b}^{[2]} = \mathbf{b}^{[2]} - \eta \frac{\partial E}{\partial \mathbf{b}^{[2]}} = \mathbf{b}^{[2]} - \eta \boldsymbol{\delta}^{[2]} = \begin{pmatrix} -0.0005 \\ -0.0005 \\ -0.0005 \end{pmatrix}$$

$$\bullet \mathbf{W}^{[3]} = \mathbf{W}^{[3]} - \eta \frac{\partial E}{\partial \mathbf{W}^{[3]}} = \mathbf{W}^{[3]} - \eta \boldsymbol{\delta}^{[3]} (\mathbf{x}^{[2]})^T = \begin{pmatrix} 0.1 & 0.1 & 0.1 \\ 0.1 & 0.1 & 0.1 \end{pmatrix} - 1 \begin{pmatrix} -0.114 \\ 0.134 \end{pmatrix} (0.54 \quad 0.54 \quad 0.54) = \begin{pmatrix} 0.162 & 0.162 & 0.162 \\ 0.027 & 0.027 & 0.027 \end{pmatrix}$$


$$\bullet \mathbf{b}^{[3]} = \mathbf{b}^{[3]} - \eta \frac{\partial E}{\partial \mathbf{b}^{[3]}} = \mathbf{b}^{[3]} - \eta \boldsymbol{\delta}^{[3]} = \begin{pmatrix} 0.114 \\ -0.134 \end{pmatrix}$$

$$\bullet E(\mathbf{w}_{init}) = -\frac{1}{2}((1 - 0.54)^2 + (0 - 0.54)^2) = 0.25, \quad E(\mathbf{w}_{new}) = -\frac{1}{2}((1 - 0.593)^2 + (0 - 0.544)^2) = 0.23$$

■ Notes:

- considering the classification of the given observation, we observe that it decreased from 0.25 to 0.23
- random weight initializations will prevent repeated values on $\boldsymbol{\delta}^{[p]}$, promoting a more expressive learning
- first layers update more slowly than last layers (we will return to this point later!)

Outline



- Multi-layer perceptron
 - non-linearity
 - gradient descent
- Propagation
- **Backpropagation**
 - network updates
 - tensor operations
 - **batch *versus* stochastic updates**
 - cross-entropy
- Learning convergence
 - optimality
 - early stopping

Gradient descent variants

- **Batch**, classic or steepest gradient descent


- use all n observations to compute $\frac{\partial E}{\partial \mathbf{W}^{[p]}}$ contributions
- sum all the contributions and only after update weights

$$\mathbf{W}^{[p]} = \mathbf{W}^{[p]} - \eta \left(\frac{\partial E^{(1)}}{\partial \mathbf{W}^{[p]}} + \dots \frac{\partial E^{(n)}}{\partial \mathbf{W}^{[p]}} \right)$$

- some disadvantages:
 - the low stochasticity can lead to be stuck in local minima
 - efficiency: updates can be slow to compute for high n

- **Stochastic** gradient descent

- use one observation at a time $\mathbf{W}^{[p]} = \mathbf{W}^{[p]} - \eta \frac{\partial E^{(i)}}{\partial \mathbf{W}^{[p]}}$
- disadvantage: weak stability can be associated with a heightened number of iterations until convergence




— Batch gradient descent
— Stochastic gradient descent


Gradient descent variants

- Can we trade-off the limitations of batch and stochastic GD?
 - Yes: mini-batch gradient descent
- A **mini-batch** is a compact subset of observations
- Mini-batches are randomly selected to provide:
 - more stable updates than stochastic alternatives
 - more efficient updates than the classic GD while more able to escape local minimum
- How to choose the #observations in a mini-batch?
- The number of **epochs** is generally used in reference to the number of (mini-batch) network updates


Gradient Descent



Stochastic Gradient Descent




Mini-Batch Gradient Descent



Mini-batch backpropagation

- Mini-batch backpropagation with batches of size q :
 1. Initialize the weights to small random values
 2. Randomly choose q observations from X , $\{\mathbf{x}_{i_1}, \dots, \mathbf{x}_{i_q}\}$
 3. Propagate the signals through the network, $\mathbf{x}^{(i)[p]} = \phi^{[p]}(\mathbf{W}^{[p]}\mathbf{x}^{(i)[p-1]} + \mathbf{b}^{[p]})$
 4. Compute the deltas for the output layer $\boldsymbol{\delta}^{[P]} = (\mathbf{x}^{[(i)P]} - \mathbf{t}^{(i)}) \circ \phi'^{[P]}(\mathbf{z}^{(i)[P]})$
 5. Compute the deltas for the preceding layers $\boldsymbol{\delta}^{[p]} = \mathbf{W}^{[p+1]T} \cdot \boldsymbol{\delta}^{(i)[p+1]} \circ \phi'^{[p]}(\mathbf{z}^{(i)[p]})$
 6. Update all connections, $\mathbf{W}^{[p]} = \mathbf{W}^{[p]} - \eta \left(\frac{\partial E}{\partial \mathbf{W}^{[p]}}^{(1)} + \dots \frac{\partial E}{\partial \mathbf{W}^{[p]}}^{(q)} \right)$, where $\frac{\partial E}{\partial \mathbf{W}^{[p]}}^{(i)} = \boldsymbol{\delta}^{(i)[p]} (\mathbf{x}^{(i)[p-1]})^T$
 7. Compute the loss $E(\mathbf{w})$
 8. **If** there is evidence of convergence (loss not improving for past iterations) **then** terminate
Else go to (2) and repeat for the next batch

Outline



- Multi-layer perceptron
 - non-linearity
 - gradient descent
- Propagation
- **Backpropagation**
 - network updates
 - tensor operations
 - batch *versus* stochastic updates
 - **cross-entropy**
- Learning convergence
 - optimality
 - early stopping

Cross-entropy

- Until here, our focus has been placed on minimizing squared error function

$$E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n \sum_{l=1}^{|C|} \left(t_l^{(i)} - o_l^{(i)} \right)^2$$

- yet **alternative loss functions** yield relevant properties
- cross-entropy is a measure of concordance, classically $H(P, Q) = - \sum_{i=1}^{|C|} p_i \log(q_i)$
where P is our ground truth distribution (*targets*) and Q is the learnt distribution (*estimates*)
- it can be used as the loss: **cross-entropy error**...

$$E(\mathbf{w}) = - \sum_{i=1}^n \sum_{l=1}^{|C|} t_l^{(i)} \log(o_l^{(i)})$$

- Exercise: given targets $\left\{ \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}, \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \right\}$ and MLP estimates $\left\{ \begin{pmatrix} 0.3 \\ 0.5 \\ 0.2 \end{pmatrix}, \begin{pmatrix} 0.8 \\ 0.1 \\ 0.1 \end{pmatrix} \right\}$ compute the MLP's cross-entropy

Cross-entropy and activations

- An important property of cross entropy:
 - outputs o_l should be in $[0,1]$. Why? Inspect the loss
 - furthermore, the classic cross-entropy formulation, $H(P, Q)$, generally assumes:
 $\sum_i p_i = 1$ and $\sum_i q_i = 1$ (in other words $\sum_l t_l = 1$ and $\sum_l o_l = 1$)
- Important implications:
 - the activation functions at the output layer, i.e. , should be also in $[0,1]$
 - e.g. $\text{sigmoid}(x)$ is eligible, $\text{ReLU}(x)$ is not eligible
 - although not mandatory, the output activations should $\sum_l o_l = 1$
 - can we simply normalize the output values $\frac{o_k}{\sum_l o_l}$?
 - not as simple! If we use sigmoid activation and post-normalize the outputs, normalization will affect the learning – need to revise the update rules!


Softmax

- At the light of previous limitations – solution?

- **softmax** activation

$$\text{softmax}(z_i^{[P]} | \mathbf{z}^{[P]}) = \frac{e^{z_i^{[P]}}}{\sum_{l=1}^L e^{z_l^{[P]}}}$$

- preferred activation on the output layer to use with cross-entropy error
- properties:
 - ensures $\sum_l o_l = 1$
 - focus on differences instead of absolute values
 - exercise: find the softmax outputs for $\mathbf{z}^{[P]} = (1 \ 4 \ 0)^T$, $\mathbf{z}^{[P]} = (11 \ 14 \ 10)^T$ and $\mathbf{z}^{[P]} = (0.25 \ 1 \ 0)^T$
 - this confers expressivity to the learning



Minimizing cross-entropy

- Let us find the update rules for cross-entropy using general activations $\phi^{[p]}$
- Knowing $o_l = \phi^{[P]}(\sum_{k=1}^K w_{kl}^{[P]} x_k^{[P-1]})$... we get

$$E(\mathbf{w}) = - \sum_{l=1}^2 \sum_{i=1}^n t_l^{(i)} \log \left(\phi^{[P]} \left(\sum_{k=1}^K w_{kl}^{[P]} v_k^{[P-1]} \right) \right)$$

- recalling that $\frac{\partial \log_a x}{\partial x} = \frac{1}{x} \ln(a)$ and considering $a = e$ for simplicity
- when minimizing the error...

$$\frac{\partial E}{\partial w_{kl}^{[P]}} = \dots \text{some steps later} \dots = - \sum_{i=1}^n \left(t_l^{(i)} - o_l^{(i)} \right) \cdot x_k^{(i)[P-1]}$$

$$\frac{\partial E}{\partial \mathbf{w}^{[P]}} = \sum_{i=1}^n \boldsymbol{\delta}^{[P]} \cdot \mathbf{x}^{(i)[P-1]T} \text{ with } \boldsymbol{\delta}^{[P]} = (\mathbf{o}^{(i)} - \mathbf{t}^{(i)})$$

Minimizing cross-entropy

- For the former layer, by applying the chain rule... and a few mathematics... we get


$$\frac{\partial E}{\partial w_{jk}^{[p]}} = \sum_{i=1}^n \frac{\partial E}{\partial x_k^{(i)[p]}} \cdot \frac{\partial x_k^{(i)[p]}}{\partial w_{jk}^{[p]}} = \dots = - \sum_{i=1}^n \sum_{l=1}^2 \left(t_l^{(i)} - o_l^{(i)} \right) \cdot w_{kl}^{(i)[p+1]} \cdot \phi'(z_k^{(i)[p]}) \cdot x_j^{(i)[p-1]}$$

- yielding the same updates for any former layer:

$$\frac{\partial E}{\partial \mathbf{W}^{[p]}} = \boldsymbol{\delta}^{[p]} \mathbf{x}^T \quad \text{with} \quad \boldsymbol{\delta}^{[p]} = \mathbf{W}^{[p+1]T} \cdot \boldsymbol{\delta}^{[p+1]} \circ \phi'^{[p]}(\mathbf{z}^{[p]})$$

- What is then the **difference** between cross-entropy and SSE?
 - just $\boldsymbol{\delta}^{[P]}$ (which then affects every other delta $\boldsymbol{\delta}^{[p]}$ for $1, 2, \dots, P-1$)
 - while for the cross-entropy error we have $\boldsymbol{\delta}^{[P]} = (\mathbf{o} - \mathbf{t})$
 - for the old squared error we have $\boldsymbol{\delta}^{[P]} = (\mathbf{o} - \mathbf{t}) \circ \phi'^{[P]}(\mathbf{z}^{[P]})$
 - this difference in red makes cross-entropy more attractive in some scenarios as it can yield faster convergence

Outline



- Multi-layer perceptron
 - non-linearity
 - gradient descent
- Propagation
- Backpropagation
 - network updates
 - tensor operations
 - batch *versus* stochastic updates
 - cross-entropy
- **Learning convergence**
 - **optimality**
 - **early stopping**

Convergence of backpropagation


- As we saw with perceptron learning, we can control η to promote convergence:
 - when η is small: smooth path (yet slow)
 - when η is large: oscillatory path
 - when η exceeds a certain critical value, backpropagation can become unstable
- Still for small η ... gradient descent can find hard to leave **local minimum**
 - Problem: perhaps not a global minimum
 - Solutions?
 - **stochastic** gradient descent
 - train multiple networks with **different initial weights** (i.e., run multiple times)
 - add **momentum** (next class!)

Validation and test data


- In neural processing, data is generally divided into:
 - **training data**: to update the network
 - **validation data**: to decide the best network (parameters)...
 - more later!
 - **test data**: to get a final unbiased assessment of the network performance
 - generally we expect testing error to be worse than validation error
- k -fold Cross-Validation (multiple train-test partitions) can be as well considered for a more comprehensive comparison of neural networks
 - for each iteration, the training partition is then subdivided onto train and validation sets
 - exercise: given 1000 observations, 10-fold CV, and 90-10 train-validation split.
How many observations are set aside for validation per CV iteration? Answer: 810!
 - CV is nevertheless less common for very large networks since efficiency is hampered k times

Convergence: early stopping

- We can keep optimizing the network weights...
 - ... until the network perfectly overfits the training data, hampering the ability to generalize to unseen data
- One possible solution? **Early stopping**
 - stop convergence before MLP overfits data
 - **how?**
 - optimize weights with training data, yet **assess the loss on the validation set**
 - stop learning when the validation error increases along few iterations (evidence of overfitting)



Outline



- **Multi-layer perceptron**
 - non-linearity
 - gradient descent
- **Propagation**
- **Backpropagation**
 - network updates
 - tensor operations
 - batch *versus* stochastic updates
 - cross-entropy
- **Learning convergence**
 - optimality
 - early stopping

Thank You



rmch@tecnico.ulisboa.pt
andreas.wichert@tecnico.ulisboa.pt